



Project Log:

LLM_RAG_Financial_Dashboard.ipynb

1 What Was Done

- ☒ Loaded a cleaned financial dataset (CSV format).
 - ☒ Created a function `human_readable_volume(v)` to format large volume numbers (e.g., `24409527` → `24,409,527`).
 - ☒ Used a loop to apply that function and created a new column: `Formatted_Volume`.
 - ☒ Created another function `create_readable_sentence(row)` to convert each row into a natural language summary.
 - ☒ Used `df.iterrows()` to loop through each row and call the sentence function.
 - ☒ Stored each sentence in a list called `text_chunks`.
 - ☒ Saved all the summary sentences into a file: `financial_text_chunks.txt` (1 row per line).
-

2 Why It Was Done

- GPT and other LLMs understand and work best with **plain natural language**, not raw tables.
- Turning structured financial data into **text summaries** allows LLMs to reason over your dataset.
- These text chunks will be used in a **RAG (Retrieval-Augmented Generation)** setup:

- The LLM will search for relevant sentences from your file before answering a user's query.
 - It builds the foundation for enabling **chat-style Q&A over your own financial data**.
-

3 Challenges Faced & How They Were Resolved

✗ Error:

`TypeError: tuple indices must be integers or slices, not str`

- **Cause:** Looping with `for row in df.iterrows()` gave a tuple, not a DataFrame row.
 - **✓ Fix:** Replaced with `for _, row in df.iterrows()` — now `row` is a proper Series.
-

✗ Issue:

Newline characters not working — all sentences appeared on one line in the text file.

- **Cause:** Used `"\\n"` instead of `"\n"` when writing to file.
 - **✓ Fix:** Corrected to `f.write(line + "\n")` to properly separate lines in `financial_text_chunks.txt`.
-

4 Why Each Step Was Important

Step	Purpose
<code>human_readable_volume(v)</code>	Adds commas to large numbers so text is easy to read
<code>Formatted_Volume</code> column	Avoids repeating formatting each time a sentence is created
<code>create_readable_sentence(row)</code>	Turns each table row into a clear sentence GPT can understand

<code>df.iterrows()</code>	Lets you access and work with each row of data one by one
<code>text_chunks</code> list	Collects all summaries in one place for export
<code>financial_text_chunks.txt</code>	This is your custom knowledge base for RAG search

5 Outcome & Current Project Status

Task	Status
Data loaded and cleaned	✓ Done
Volume formatted	✓ Done
Sentences generated	✓ Done
Text chunks saved to file	✓ Done
Ready for embeddings	✓ YES ✓

Next Step:

- Generate **embeddings** from `financial_text_chunks.txt`
- Store in **vector DB** (FAISS or ChromaDB)
- Connect with **LLM using LangChain or LlamaIndex**
- Build your **Streamlit interface for querying the data**